

SVKM'S NMIMS (Deemed-to-be University)
MUKESH PATEL SCHOOL OF TECHNOLOGY MANAGEMENT AND ENGINEERING
NAVI MUMBAI CAMPUS

Roll Number: A053	Name: Nishant Baruah
Class: B. Tech CE (Sem VII)	Batch: B2
Date of Experiment: 22-10-24	Date of Submission: 05-11-24
Subject: Cloud Computing	Practical No: 10

AIM:

Lab Manual: Working with Kubernetes to run a highly available web server

Task:

```
This is a sandbox environment. Using personal credentials
is HIGHLY discouraged. Any consequences of doing so, are
completely the user's responsibilities.

You can bootstrap a cluster as follows:

1. Initializes cluster master node:

kubeadm init --apiserver-advertise-address $(hostname -i) --pod-network-cidr 10.5.0.0/16

2. Initialize cluster networking:

kubectrl apply -f https://raw.githubusercontent.com/cloudnativelabs/kube-router/master/daemonset/kubeadm-kuberouter.yaml

3. (Optional) Create an nginx deployment:

kubectrl apply -f https://raw.githubusercontent.com/kubernetes/website/master/content/en/examples/application/nginx-app.yaml

The FWK team.

[node1 ~]$ kubeadm init --apiserver-advertise-address $(hostname -i) --pod-network-cidr 10.5.0.0/16
Initializing machine ID from random generator.
#1023 19:31:46.809001 484 initconfiguration.go:120] Usage of CRI endpoints without URL scheme is deprecated and can cause kubelet errors in the futu
re. Automatically prepending scheme "unix" to the "criSocket" with value "/run/docker/containerd/containerd.sock". Please update your configuration!
#1023 19:31:46.994343 484 version.go:256] remote version is much newer: v1.31.2; falling back to: stable-1.27
[init] Using Kubernetes version: v1.27.16
[preflight] Running pre-flight checks
[preflight] The system verification failed. Printing the output from the verification:
KERNEL_VERSION: 4.1.0-210-generic
OS: Linux
CGROUPS_CPU: enabled
CGROUPS_CPUACCT: enabled
CGROUPS_CPUSET: enabled
CGROUPS_DEVICES: enabled
CGROUPS_FREEZER: enabled
CGROUPS_MEMORY: enabled
CGROUPS_PIDS: enabled
CGROUPS_HUGETLB: enabled
CGROUPS_BLKIO: enabled
[WARNING SystemVerification]: failed to parse kernel config: unable to load kernel module: "configs", output: "", err: exit status 1
[WARNING FileContent--proc-sys-net-bridge-bridge-nf-call-iptables]: /proc/sys/net/bridge/bridge-nf-call-iptables does not exist
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
#1023 19:31:47.530833 484 images.go:80] could not find officially supported version of etcd for Kubernetes v1.27.16, falling back to the nearest etc
d version (3.5.7-0)
```

```
[wait-control-plane] Waiting for the kubelet to boot up the control plane as static Pods from directory "/etc/kubernetes/manifests". This can take up to 4m0s
[apiclient] All control plane components are healthy after 6.004001 seconds
[upload-config] Storing the configuration used in ConfigMap "kubeadm-config" in the "kube-system" Namespace
[kubelet] Creating a ConfigMap "kubelet-config" in namespace kube-system with the configuration for the kubelets in the cluster
[upload-certs] Skipping phase. Please see --upload-certs
[mark-control-plane] Marking the node node1 as control-plane by adding the labels: [node-role.kubernetes.io/control-plane node.kubernetes.io/exclude-from-external-load-balancers]
[mark-control-plane] Marking the node node1 as control-plane by adding the taints [node-role.kubernetes.io/control-plane:NoSchedule]
[bootstrap-token] Using token: emj87o.cwrix5r5y6jubzhz
[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 192.168.0.23:6443 --token emj87o.cwrix5r5y6jubzhz \
--discovery-token-ca-cert-hash sha256:57a04cdeb6fa4265e9a98c89ea380578b1fca096633ecdef0163b399bbe2057c0
```

Waiting for api server to start up

Warning: resource daemonset/kube-proxy is missing the kubect1.kubernetes.io/last-applied-configuration annotation which is required by kubect1 apply. kubect1 apply should only be used on resources created declaratively by either kubect1 create --save-config or kubect1 apply. The missing annotation will be patched automatically.

daemonset.apps/kube-proxy configured

No resources found

```
[node1 ~]$ kubect1 apply -f https://raw.githubusercontent.com/cloudnativelabs/kube-router/master/daemonset/kubeadm-kuberouter.yaml
configmap/kube-router-cfg created
daemonset.apps/kube-router created
serviceaccount/kube-router created
clusterrole.rbac.authorization.k8s.io/kube-router created
clusterrolebinding.rbac.authorization.k8s.io/kube-router created
[node1 ~]$
```

```
[node2 ~]$ kubeadm join 192.168.0.23:6443 --token 8lrstv.jno0j886vt6sx2vu --discovery-token-ca-cert-hash sha256:57a04cdeb6fa4265e9a98c89ea380578b1fca096633ecdef0163b399bbe2057c0
```

Initializing machine ID from random generator.

W1023 19:35:47.911619 865 initconfiguration.go:120] Usage of CRI endpoints without URL scheme is deprecated and can cause kubelet errors in the future. Automatically prepending scheme "unix" to the "criSocket" with value "/run/docker/containerd/containerd.sock". Please update your configuration!

[preflight] Running pre-flight checks

[preflight] The system verification failed. Printing the output from the verification:

KERNEL_VERSION: 4.4.0-210-generic

OS: Linux

CGROUPS_CPU: enabled

CGROUPS_CPUACCT: enabled

CGROUPS_CPUSET: enabled

CGROUPS_DEVICES: enabled

CGROUPS_FREEZER: enabled

CGROUPS_MEMORY: enabled

CGROUPS_PIDS: enabled

CGROUPS_HUGEBLB: enabled

CGROUPS_BLKIO: enabled

[WARNING SystemVerification]: failed to parse kernel config: unable to load kernel module: "configs", output: "", err: exit status 1

[WARNING FileContent--proc-sys-net-bridge-bridge-nf-call-iptables]: /proc/sys/net/bridge/bridge-nf-call-iptables does not exist

[preflight] Reading configuration from the cluster...

[preflight] FYI: You can look at this config file with 'kubect1 -n kube-system get cm kubeadm-config -o yaml'

[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"

[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"

[kubelet-start] Starting the kubelet

[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:

* Certificate signing request was sent to apiservert and a response was received.

* The Kubelet was informed of the new secure connection details.

```
[node1 ~]$ kubect1 get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
node1	Ready	control-plane	3m55s	v1.27.2
node2	Ready	<none>	18s	v1.27.2

```
[node1 project]$ vi deployment.yaml
```

```
[node1 project]$ kubect1 apply -f deployment.yaml
```

deployment.apps/nginx-deployment created

```
[node1 project]$ kubect1 get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-584dbbbc9f-5r88j	1/1	Running	0	13s
nginx-deployment-584dbbbc9f-69xww	1/1	Running	0	13s
nginx-deployment-584dbbbc9f-g79rz	1/1	Running	0	13s

```
[node1 project]$
```

Conclusion

Learned how to successfully set up a Nginx web server running with 3 replicas using Kubernetes cluster.